

Adding a positive/negative cache to the term comparison routine

Enrico Tassi

INRIA Saclay

14 Feb 2013 — Paris

Roadmap

The problem

The idea

Work in progress

Some results

Long story short

Computation in CIC (reduction):

- ▶ plays an important role in CIC
- ▶ optimized for > 10 years (call by need, *_compute)

Type checking compares terms up-to reduction (conversion):

big scale (e.g. reflection)

small scale (e.g. regular type checking)

As of today:

- ▶ the big case is mainly reduction, we are good!
- ▶ very few optimizations for the “small” one...

The current algorithm

Conversion test (reduction.ml):

```
let rec are_conv t1' t2' =  
  let t1 = whd_nodelta t1' and t2 = whd_nodelta t2' in  
  if are_App_with_same_head t1 t2 then  
    try List.forall2 are_conv (args t1) (args t2)  
    with Err -> try are_conv t1 (unfold_head t2)  
    with NotUnfold -> are_conv (unfold_head t1) t2  
  else ...
```

Call by need (as in closure.ml):

- ▶ employs mutable terms (fconstr/fterm)
- ▶ weak head normal forms are “cached” in place

So what?

What's wrong with this approach:

- ▶ when a term has variables its reducts are usually bigger
- ▶ hence sharing reductions means sharing expansions
- ▶ comparing expanded terms is *way* more costly

Not so artificial example:

- ▶ testing: $\text{proj_op } S1 \ 0 \ 10! \equiv \text{proj_op } S2 \ 10 \ 9!$
may mean: $0 + 10! \equiv 10 * 9!$
as well as: $0 + 10! \equiv 10 + 9!$

Pitfalls

We do use HO constructs and *today* we have big terms even while proving:

- ▶ HO may trick conversion to normalize sub-terms first
- ▶ every comparison that follows works with bigger terms (imperative updates of the CBN machine)

Work around

In the Mathematical Components library:

- ▶ locking of some transparent definitions
- ▶ inline the definition of some constants (projection) to force eager head reduction
- ▶ ...

In the (old) Univalent Foundations library:

- ▶ patch `closure.ml` to implement call by name

Roadmap

The problem

The idea

Work in progress

Some results

Congruence closure

Used when terms have to be tested for equality

- ▶ using transitivity ($a = b \rightarrow c = b \rightarrow a = c$)
- ▶ using congruence ($a = b \rightarrow f\ a = f\ b$)

In the context of SMT solvers, ML type checking, etc.

With a bit of sugar it resembles conversion:

- ▶ $a \equiv c$ if $a \triangleright b$ and $c \triangleright b$
- ▶ $f\ a \equiv f\ b$ if $a \equiv b$ (otherwise we unfold f and continue)

Admits efficient implementation (based on union-find).

A negative/positive cache

Union-find requires:

- ▶ fast syntactic comparison of terms (usually an int)
- ▶ fast associative map over terms (array: `int -> int`)

Could be read as: fast hash tables on “terms”, that requires:

- ▶ fast comparison
- ▶ fast hash calculation

Roadmap

The problem

The idea

Work in progress

Some results

My approach

Steps:

1. dump all conversion problems (OK)
2. hashconsing for term and red machine (ok)
3. reduction machine on new data structures (ok)
4. union-find cache (almost)
5. testing (little)
6. tune conversion strategy (todo)

Where: [github.com/gares/coq/ branch speedup/tcomp](https://github.com/gares/coq/branch/speedup/tcomp)

Opportunistic hashconsing

Fast hash tables means:

- ▶ fast comparison (hashconsing)
- ▶ fast hashing (cache the hash value)

The idea

- ▶ non-leaf nodes have an extra `int` field
- ▶ `mkBla` sets it to 0 (invalid hash value)
- ▶ internalizing a term updates the hash value
- ▶ invariant: non zero hash field means already canonical

Advantage: easy/progressive integration

New conversion test

Conversion as in conversion.ml:

```
let rec are_conv t1' t2' =
  let t1 = intern (whd_nodelta t1') in UF.union t1 t1';
  let t2 = intern (whd_nodelta t2') in UF.union t2 t2';
  match UF.same t1 t2 with
  | 'No -> raise Err
  | 'Yes -> ()
  | 'Maybe ->
    if are_App_with_same_head t1 t2 then
      try List.forall2 are_conv (args t1) (args t2);
        UF.union t1 t2
      with Err -> ...
    else
      ... UF.partition t1 t2; raise Err ...
```

Roadmap

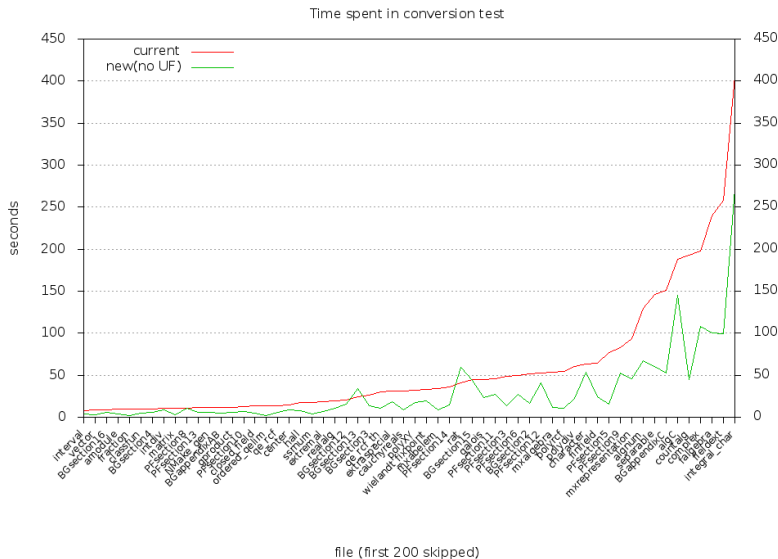
The problem

The idea

Work in progress

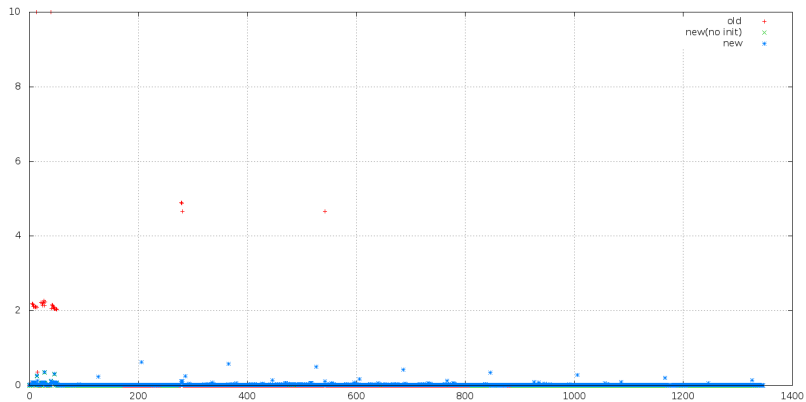
Some results

The conversion machine w.o. the cache

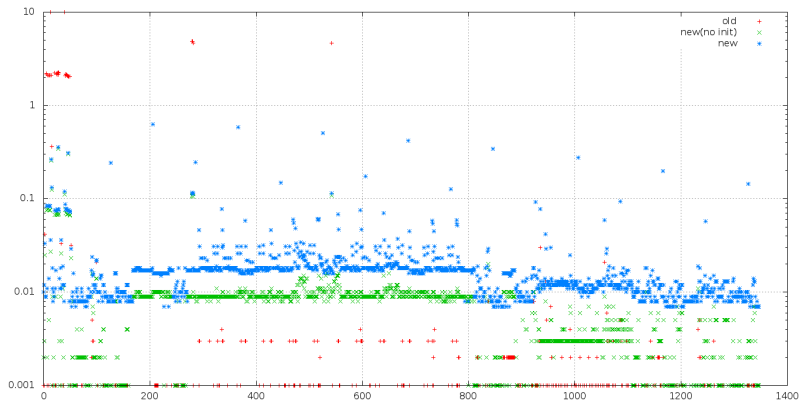


Is there a bench for pure reduction?

The conversion machine on uu0



The conversion machine on uu0 (log scale)



Open questions

- ▶ Negative cache (now very naive)
- ▶ Better hash tables (faster allocation or faster resizing)
- ▶ Better use of opportunistic sharing outside the kernel (unification piggy backs on the kernel)